

Training pipeline

1 · Generate
BODHI CoT on
MedGemma-27B
4,000 prompts, greedy

2 · Grade
Qwen-14B rubric
per-criterion
score $s \in [0, 1]$

3 · Filter
Retain $s \geq 0.4$
90/10 train-val
 $\sim 3,100$ traces/seed

4 · Distill
LoRA SFT
 $r=16$, 3 epochs
5 seeds

Base model
MedGemma-27B

Base
no CoT

Base + BODHI CoT
CoT at inference

LoRA adapter
(5 seeds)

LoRA
no CoT

LoRA + BODHI CoT
CoT at inference

HealthBench-Hard (held-out)
1,000 prompts · Llama-3.1-8B-Instruct grader

Evaluation (2×2 factorial, out-of-sample)