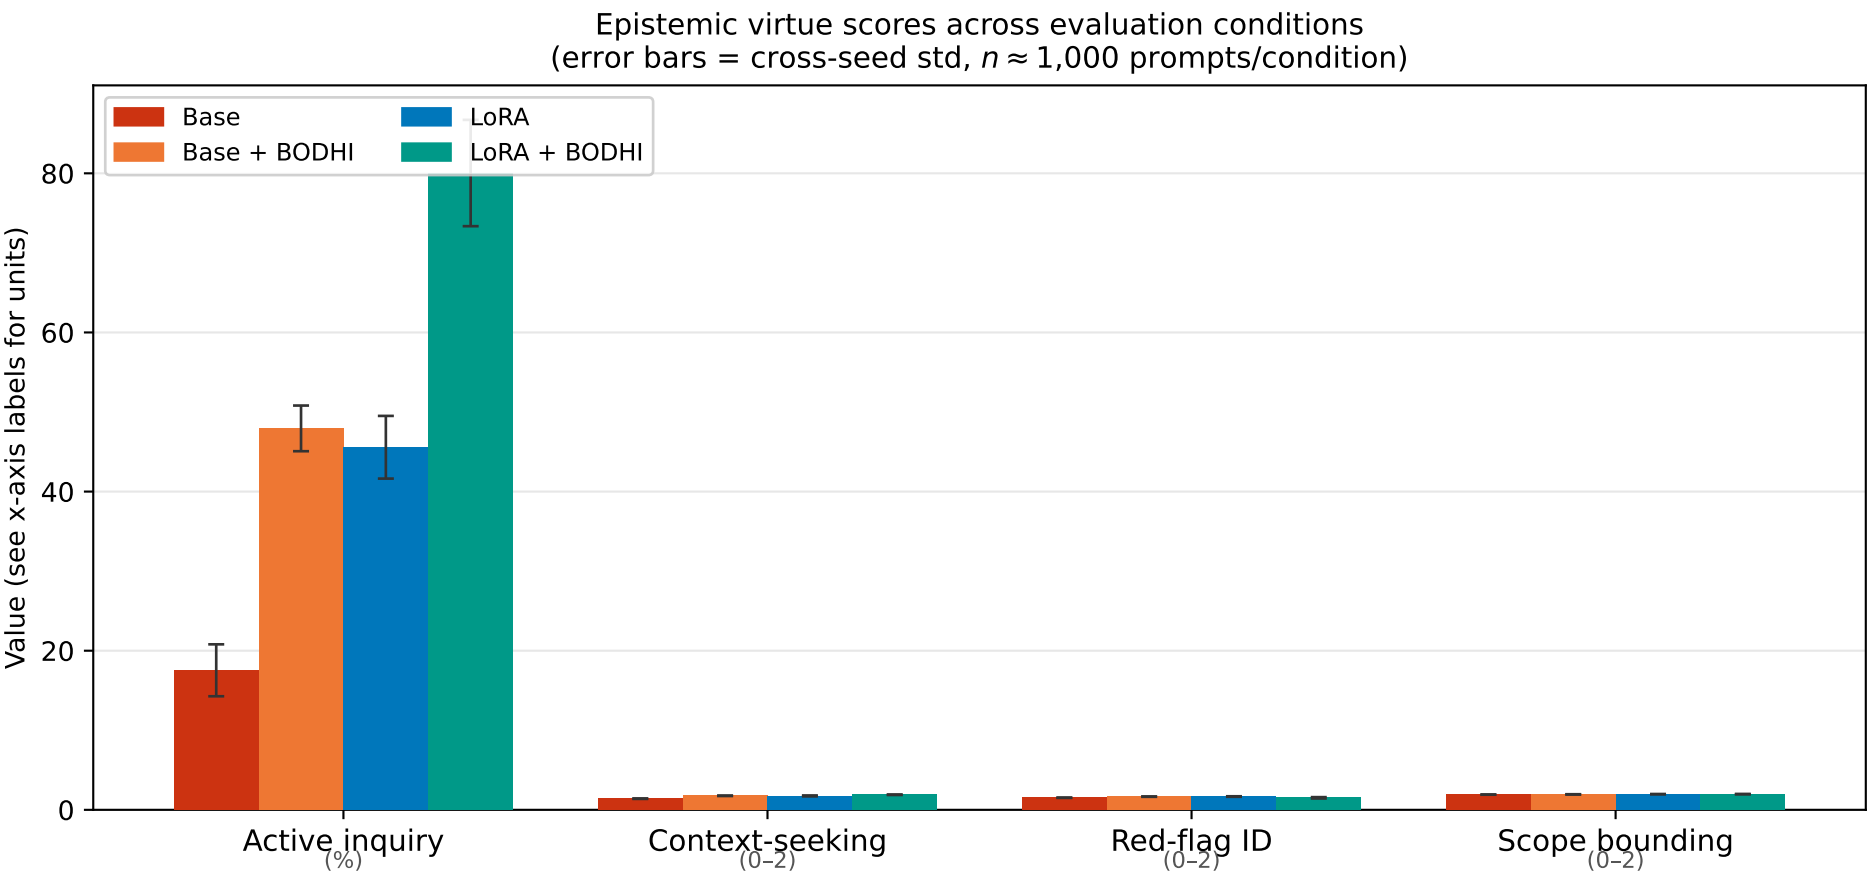
Epistemic virtue scores across evaluation conditions
(error bars = cross-seed std, n ≈ 1,000 prompts/condition)
Base
Base + BODHI
LoRA
LoRA + BODHI
Value (see x-axis labels for units)
0
20
40
60
80
Active inquiry
(%)
Context-seeking
(0–2)
Red-flag ID
(0–2)
Scope bounding
(0–2)